Newton update in L_2-norm random tree approximation

Hongyu Su

Helsinki Institute for Information Technology HIIT
Department of Computer Science
Aalto University

May 20, 2015

Preliminaries

- $\mathcal{X}$ is an arbitrary input space, $\mathbf{x} \in \boldsymbol{\mathcal{X}}$.
- $\mathcal{Y}$ is an output space of a set of ℓ-dimensional *multilabels*

$$\mathbf{y} = (y_1, \cdots, y_\ell) \in \boldsymbol{\mathcal{Y}}.$$

- y_i is a *microlabel* and $y_i \in \{1, \cdots, r_i\}, r_i \in \mathbb{Z}$.
- For example, multilabel binary classification $y_i \in \{-1, +1\}$.
- Training examples are sampled from $(\mathbf{x}, \mathbf{y}) \in \boldsymbol{\mathcal{X}} \times \boldsymbol{\mathcal{Y}}$.
- Each example $(\mathbf{x}, \mathbf{y})$ is mapped into a joint feature space $\boldsymbol{\phi}(\mathbf{x}, \mathbf{y})$.
- $\mathbf{w}$ is the weight vector in the joint feature space.
- Define a linear score function $F(\mathbf{w}, \mathbf{x}, \mathbf{y}) = \langle \mathbf{w}, \boldsymbol{\phi}(\mathbf{x}, \mathbf{y}) \rangle$.
- The prediction $\mathbf{y_w}(\mathbf{x})$ of an input $\mathbf{x}$ is the multilabel $\mathbf{y}$ that maximizes the score function

$$\mathbf{y_w}(\mathbf{x}) = \underset{\mathbf{y} \in \boldsymbol{\mathcal{Y}}}{\mathbf{argmax}} \, \langle \mathbf{w}, \boldsymbol{\phi}(\mathbf{x}, \mathbf{y}) \rangle. \tag{1}$$

- (??) is called *inference* problem which is $\mathcal{NP}$-hard for most output feature maps.

Markov network

- We assume that the output feature map ϕ is a potential function on a Markov network $G = (E, V)$.

- G is a complete graph with $|V| = \ell$ nodes and $|E| = \frac{\ell(\ell-1)}{2}$ undirected edges.

- $\varphi(\mathbf{x})$ is the input feature map, e.g., bag-of-words feature of an example $\mathbf{x}$.

- $\psi(\mathbf{y})$ is the output feature map which is a collection of edges and labels

$$\varphi(\mathbf{y}) = (u_e)_{e \in E}, u_e \in \{-1, +1\}^2.$$

- The joint feature is the Kronecker product of $\varphi(\mathbf{x})$ and $\psi(\mathbf{y})$

$$\phi(\mathbf{x}, \mathbf{y}) = (\phi_e(\mathbf{x}, \mathbf{y}))_{e \in E} = (\varphi(\mathbf{x}) \otimes \psi_e(\mathbf{y}_e))_{e \in E}.$$

- The score function is

$$F(\mathbf{w}, \mathbf{x}, \mathbf{y}) = \langle \mathbf{w}, \phi(\mathbf{x}, \mathbf{y}) \rangle = \sum_{e \in E} \langle \mathbf{w}_e, \phi_e(\mathbf{x}, \mathbf{y}_e) \rangle.$$

- Solving the following inference on a complete graph is difficult

$$\mathbf{y_w}(\mathbf{x}) = \operatorname*{argmax}_{\mathbf{y} \in \mathcal{Y}} \sum_{e \in E} \langle \mathbf{w}_e, \boldsymbol{\phi}_e(\mathbf{x}, \mathbf{y}_e) \rangle.$$

Multilabel classification

- Multilabel classification is an important research field in machine learning.

    - For example, a document can be classified as "science", "genomics", and "drug discovery".
    - Each input variable $\mathbf{x} \in \boldsymbol{\mathcal{X}}$ is associated with multiple output variables $\mathbf{y} \in \boldsymbol{\mathcal{Y}}, \boldsymbol{\mathcal{Y}} = \mathcal{Y}_1 \times \cdots \times \mathcal{Y}_l, \mathcal{Y}_i = \{+1, -1\}$.
    - The goal is to find a mapping function that predicts the best values of an output given an input $f \in \mathcal{H} : \boldsymbol{\mathcal{X}} \to \boldsymbol{\mathcal{Y}}$.

- The central problems of multilabel classification:

    - The size of the output space $\boldsymbol{\mathcal{Y}}$ is exponential in the number of microlabels.
    - The dependency of microlabels needs to be exploited to improve the prediction performance.

Structured output learning

- There is an *output graph* connecting multiple labels.
  - A set of nodes represents multiple labels.
  - A set of edges represents the correlation between labels.
- Hierarchical classification:
  - The output graph is a rooted tree or a directed graph defining different levels of granularities.
  - For example, SSVM, ...
- Graph labeling:
  - The output graph often takes a general form (e.g., a tree, a chain).
  - For example, M^3N, CRF, MMCRF, ...
- The output graph is assumed to be known *apriori*.

Research question

- The output graph is hidden in many applications.
    - For example, a surveillance photo can be tagged with "building", "road", "pedestrian", and "vehicle".
- We study the problem in structured output learning when the output graph is not observed.
- In particular:
    - Assume the dependency can be expressed by a complete set of pairwise correlations.
    - Build a structured output learning model with a complete graph as the output graph.
    - Solve the optimization problem and the inference problem ($\mathcal{NP}$-hard).

Today

- A structured prediction model which performs max-margin learning on a random sample of spanning tree.

- Two ways to combine the set of random spanning trees
  - conical combination in NIPS paper.
  - convex combination as future work.

- Derivations and the corresponding optimization problems.

Model

- Training examples comes in pair $S = \{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^m \in \mathcal{X} \times \mathcal{Y}$.

- A complete graph $G = (E, V)$ is used as the output graph.

- $\varphi(\mathbf{x})$ is the input feature map, e.g., a feature vector of d dimension.

- $\Gamma_G(\mathbf{y})$ is the output feature map of $\mathbf{y}$ on G of $4 \times |E|$ dimension

$$\Gamma_G(\mathbf{y}) = \{\Gamma_e(\mathbf{y}_e)\}_{e \in G},$$
$$\Gamma_e(\mathbf{y}_e) = [\mathbf{1}_{\mathbf{y}_e==00}, \mathbf{1}_{\mathbf{y}_e==01}, \mathbf{1}_{\mathbf{y}_e==10}, \mathbf{1}_{\mathbf{y}_e==11}]$$

- A joint feature map of $(\mathbf{x}_i, \mathbf{y}_i)$

$$\phi_G(\mathbf{x}_i, \mathbf{y}_i) = \varphi(\mathbf{x}_i) \otimes \Gamma_G(\mathbf{y}_i) = \{\phi_e(x_i, \mathbf{y}_{i,e})\}_{e \in G}.$$

- A compatibility score is defined as

$$F(\mathbf{x}, \mathbf{y}; \mathbf{w}_G) = \langle \mathbf{w}_G, \phi_G(\mathbf{x}, \mathbf{y}) \rangle = \sum_{e \in G} \langle \mathbf{w}_{G,e}, \phi_e(\mathbf{x}, \mathbf{y}_e) \rangle$$

Model (cont.)

▶ $\mathbf{w}$ ensures an input $\mathbf{x}_i$ with a correct multilabel $\mathbf{y}_i$ achieves a higher score than with any incorrect multilabel $\mathbf{y} \in \mathcal{Y}$.

▶ The predicted output $\mathbf{y}(\mathbf{x})$ for a given input $\mathbf{x}$ is computed by

$$\mathbf{y}(\mathbf{x}) = \operatorname*{argmax}_{\mathbf{y} \in \mathcal{Y}} F(\mathbf{x}, \mathbf{y}; \mathbf{w}_G) = \operatorname*{argmax}_{\mathbf{y} \in \mathcal{Y}} \sum_{e \in G} \langle \mathbf{w}_{G,e}, \phi_{G,e}(\mathbf{x}, \mathbf{y}_e) \rangle,$$

which is called *inference problem*.

▶ The inference problem is $\mathcal{NP}$-hard for most joint feature maps on the complete graph.

How to learn w on a complete graph?

- The *margin* of an example $\mathbf{x}_i$ is

$$\gamma_G(\mathbf{x}_i; \mathbf{w}_G) = F(\mathbf{x}_i, \mathbf{y}_i; \mathbf{w}_G) - \max_{\mathbf{y} \in \mathcal{Y}/\mathbf{y}_i} F(\mathbf{x}_i, \mathbf{y}; \mathbf{w}_G).$$

- $\mathbf{w}$ is solved by *max-margin principle* which aims to maximize $\gamma(\mathbf{x}_i; \mathbf{w}_G)$ over all training example $\mathbf{x}_i, i \in \{1, \cdots, m\}$.

- The inference problem on a complete graph is $\mathcal{NP}$-hardness.

- The parameter space is quadratic in the number of microlabels k.

- We aim to use a joint feature map that allows the inference problem be solved in polynomial time.

Superposition of random trees

- $S(G)$ is a complete set of spanning tree generate from G, $|S(G)| = \ell^{\ell-2}$.
- Recall
  $\phi_G(\mathbf{x}, \mathbf{y}) = \{\phi_{G,e}(\mathbf{x}, \mathbf{y}_e)\}_{e \in G}, \mathbf{w}_G = \{\mathbf{w}_{G,e}\}_{e \in G}, ||\phi_G(\mathbf{x}, \mathbf{y})|| = ||\mathbf{w}_G|| = 1.$
- $\phi_T(\mathbf{x}, \mathbf{y}) = \{\phi_e(\mathbf{x}, \mathbf{y})\}_{e \in T}$ is the projection of $\phi_G(\mathbf{x}, \mathbf{y})$ on $T \in S(G)$.
- $\mathbf{w}_T = \{\mathbf{w}_{G,e}\}_{e \in T}$ is the projection of $\mathbf{w}_G$ on $T \in S(G)$.
- Rewrite

$$F(\mathbf{x}, \mathbf{y}, \mathbf{w}_G) = \sum_{e \in G} \langle \mathbf{w}_{G,e}, \phi_{G,e}(\mathbf{x}, \mathbf{y}_e) \rangle$$

$$= \frac{1}{\ell^{\ell-2}} \sum_{T \in S(G)} \sqrt{\frac{\ell}{2}} \langle \mathbf{w}_T, \phi_T(\mathbf{x}, \mathbf{y}_e) \rangle$$

$$= \frac{1}{n} \sum_{i=1}^{n} a_{T_i} \langle \hat{\mathbf{w}}_{T_i}, \hat{\phi}_{T_i}(\mathbf{x}, \mathbf{y}_e) \rangle,$$

$$||\hat{\phi}_T(\mathbf{x}, \mathbf{y})|| = ||\hat{\mathbf{w}}_T|| = 1, \frac{1}{n} \sum_{i=1}^{n} a_{T_i}^2 = 1, \frac{1}{n} \sum_{i=1}^{n} a_{T_i} \leq 1, a_{T_i} \geq 0, n = \ell^{\ell-2}.$$

How many trees?

▶ If there is a predictor $\mathbf{w}_G$ on complete graph achieves a margin on some training data, with high probability we need n spanning tree predictors $\{\mathbf{w}_{T_i}\}_{i=1}^{n}$ to achieve a close margin. n is quadratic in terms of ℓ.

▶ Recall

$$F(\mathbf{x}, \mathbf{y}, \mathbf{w}_{\mathcal{T}}) = \frac{1}{n} \sum_{i=1}^{n} a_{T_i} \underbrace{\langle \hat{\mathbf{w}}_{T_i}, \hat{\phi}_{T_i}(\mathbf{x}, \mathbf{y}_e) \rangle}_{F(\mathbf{x}, \mathbf{y}, \mathbf{w}_{T_i})},$$

$$||\hat{\phi}_{\mathcal{T}}(\mathbf{x}, \mathbf{y})|| = ||\hat{\mathbf{w}}_{\mathcal{T}}|| = 1, \; \frac{1}{n} \sum_{i=1}^{n} a_{T_i}^2 = 1, \; \frac{1}{n} \sum_{i=1}^{n} a_{T_i} \leq 1, \; a_{T_i} \geq 0, \; \cancel{n = \ell^{\ell-2}}.$$

Conical combination

- A sample $\mathcal{T} = \{T_1, \cdots, T_n\}$ of n spanning trees drawn from G.

- Normalized feature vectors $\hat{\phi}_{T_i}(\mathbf{x}, \mathbf{y}) = \frac{\phi_{T_i}(\mathbf{x}, \mathbf{y})}{||\phi_{T_i}(\mathbf{x}, \mathbf{y})||}$, $T_i \in \mathcal{T}$.

- Normalized feature weights $\hat{\mathbf{w}}_{T_i} = \frac{\mathbf{w}_{T_i}}{||\mathbf{w}_{T_i}||}$, $T_i \in \mathcal{T}$.

- Conical combination of spanning trees

$$F(\mathbf{x}, \mathbf{y}, \mathbf{w}_{\mathcal{T}}) = \frac{1}{\sqrt{n}} \sum_{i=1}^{n} q_i \underbrace{\langle \hat{\mathbf{w}}_{T_i}, \hat{\phi}_{T_i}(\mathbf{x}, \mathbf{y}) \rangle}_{F(\mathbf{x}, \mathbf{y}, \mathbf{w}_{T_i})}$$

$$\sum_{i=1}^{n} q_i^2 = 1, \; q_i \geq 0, \; \forall i \in \{1, \cdots, n\}.$$

Conical combination (cont.)

▶ To solve $\{\mathbf{w}_{T_i}\}_{T_i \in \mathcal{T}}$, we need to work on the optimization problem

$$\min_{\xi, \gamma, \mathbf{q}, \mathcal{W}} \quad \frac{1}{2\gamma^2} + \frac{C}{\gamma} \sum_{k=1}^{m} \xi_k$$

$$\text{s.t.} \quad \frac{1}{\sqrt{n}} \sum_{i=1}^{n} q_i \langle \hat{\mathbf{w}}_{T_i}, \hat{\phi}_{T_i}(\mathbf{x}_k, \mathbf{y}_k) \rangle - \max_{\mathbf{y} \in \mathcal{Y}} \frac{1}{\sqrt{n}} \sum_{i=1}^{n} q_i \langle \hat{\mathbf{w}}_{T_i}, \hat{\phi}_{T_i}(\mathbf{x}_k, \mathbf{y}) \rangle$$

$$\geq \gamma - \xi_k, \xi_k \geq 0, \forall k \in \{1, \cdots, m\}, \sum_{i=1}^{n} q_i^2 = 1, q_i \geq 0, \forall i \in \{1, \cdots, n\}.$$

▶ This is equivalent to

$$\min_{\mathbf{w}_{T_i}, \xi_i} \quad \frac{1}{2} \sum_{i=1}^{n} \|\mathbf{w}_{T_i}\|^2 + C \sum_{k=1}^{m} \xi_k$$

$$\text{s.t.} \quad \frac{1}{\sqrt{n}} \sum_{i=1}^{n} \langle \mathbf{w}_{T_i}, \phi_{T_t}(\mathbf{x}_k, \mathbf{y}_k) \rangle - \max_{\mathbf{y} \neq \mathbf{y}_k} \frac{1}{\sqrt{n}} \sum_{i=1}^{n} \langle \mathbf{w}_{T_t}, \phi_{T_i}(\mathbf{x}_k, \mathbf{y}) \rangle \geq 1 - \xi_k,$$

$$\xi_k \geq 0, \forall \ k \in \{1, \ldots, m\}.$$

Inference Problem

▶ The inference problem of RTA is defined as finding the multilabel $\mathbf{y}_{\mathcal{T}}(\mathbf{x})$ that maximizes the sum of scores over a collection of trees

$$\mathbf{y}_{\mathcal{T}}(\mathbf{x}) = \operatorname*{argmax}_{\mathbf{y} \in \mathcal{Y}} F_{\mathcal{T}}(\mathbf{x}, \mathbf{y}; \mathbf{w}_{\mathcal{T}}) = \operatorname*{argmax}_{\mathbf{y} \in \mathcal{Y}} \sum_{t=1}^{n} \langle \mathbf{w}_{T_t}, \phi_{T_t}(\mathbf{x}, \mathbf{y}) \rangle.$$

▶ The inference problem on each individual spanning tree can be solve efficiently in $\Theta(l)$ by *dynamic programming*

$$\mathbf{y}_{T_t}(\mathbf{x}) = \operatorname*{argmax}_{\mathbf{y} \in \mathcal{Y}} F_{T_t}(\mathbf{x}, \mathbf{y}; \mathbf{w}_{T_t}) = \operatorname*{argmax}_{\mathbf{y} \in \mathcal{Y}} \langle \mathbf{w}_{T_t}, \phi_{T_t}(\mathbf{x}, \mathbf{y}) \rangle.$$

▶ There is no guarantee that there exists a tree $T_t \in \mathcal{T}$ in which the maximizer of F_{T_t} is the maximizer of $F_{\mathcal{T}}$.

Fast Inference Over a Collection of Trees

▶ For each tree T_t, instead of computing the best multilabel $\mathbf{y}_{T_t}$, we compute K-best multilabels in $\Theta(Kl)$ time

$$\mathcal{Y}_{T_t,K} = \{\mathbf{y}_{T_t,1}, \cdots, \mathbf{y}_{T_t,K}\}.$$

▶ Performing the same computation on all trees gives a candidate list of $n \times K$ multilabels in $\Theta(nKl)$ time

$$\mathcal{Y}_{T,K} = \mathcal{Y}_{T_1,K} \cup \cdots \mathcal{Y}_{T_n,K}.$$

▶ For now, we assume the best scoring multilabel of a collection of trees exists in the list $\mathcal{Y}_{T,K}$.

▶ We proved that with a high probability $\mathbf{y}_T$ will appear in $\mathcal{Y}_{T,K}$.

▶ We can identify $\mathbf{y}_T$ from $\mathcal{Y}_{T,K}$.

Convex combination

▶ A sample $\mathcal{T}$ of n spanning trees drawn from G.

▶ Normalized feature weights $\hat{\mathbf{w}}_{T_i} = \frac{\mathbf{w}_{T_i}}{||\mathbf{w}_{T_i}||}$, $T_i \in \mathcal{T}$.

▶ Normalized feature vectors $\hat{\phi}_{T_i}(\mathbf{x}, \mathbf{y}) = \frac{\phi_{T_i}(\mathbf{x},\mathbf{y})}{||\phi_{T_i}(\mathbf{x},\mathbf{y})||}$, $T_i \in \mathcal{T}$.

▶ Convex combination of spanning trees

$$F(\mathbf{x}, \mathbf{y}, \mathbf{w}_{\mathcal{T}}) = \frac{1}{n} \sum_{i=1}^{n} q_i \langle \hat{\mathbf{w}}_{T_i}, \hat{\phi}_{T_i}(\mathbf{x}, \mathbf{y}) \rangle$$

$$\sum_{i=1}^{n} q_i = 1, \ q_i \geq 0, \ \forall i \in \{1, \cdots, n\}.$$

Convex combination (cont.)

▶ To solve $\{\mathbf{w}_{T_i}\}_{T_i \in \mathcal{T}}$, we need to work on the optimization problem

$$\min_{\xi,\gamma,\mathbf{q},\mathcal{W}} \quad \frac{1}{2\gamma^2} + \frac{C}{\gamma} \sum_{k=1}^{m} \xi_k$$

$$\text{s.t.} \quad \frac{1}{n} \sum_{i=1}^{n} q_i \langle \hat{\mathbf{w}}_{T_i}, \hat{\phi}_{T_i}(\mathbf{x}_k, \mathbf{y}_k) \rangle - \max_{\mathbf{y} \in \mathcal{Y}} \frac{1}{n} \sum_{i=1}^{n} q_i \langle \hat{\mathbf{w}}_{T_i}, \hat{\phi}_{T_i}(\mathbf{x}_k, \mathbf{y}) \rangle$$

$$\geq \gamma - \xi_k, \xi_k \geq 0, \forall k \in \{1, \cdots, m\}, \sum_{i=1}^{n} q_i = 1, q_i \geq 0, \forall i \in \{1, \cdots, n\}.$$

▶ This is equivalent to

$$\min_{\mathbf{w}_{T_i}, \xi_i} \quad \frac{1}{2} \left(\sum_{i=1}^{n} ||\mathbf{w}_{T_i}|| \right)^2 + C \sum_{k=1}^{m} \xi_k$$

$$\text{s.t.} \quad \frac{1}{n} \sum_{i=1}^{n} \langle \mathbf{w}_{T_i}, \phi_{T_t}(\mathbf{x}_k, \mathbf{y}_k) \rangle - \max_{\mathbf{y} \neq \mathbf{y}_k} \frac{1}{n} \sum_{i=1}^{n} \langle \mathbf{w}_{T_t}, \phi_{T_i}(\mathbf{x}_k, \mathbf{y}) \rangle \geq 1 - \xi_k,$$

$$\xi_k \geq 0, \forall k \in \{1, \ldots, m\}.$$

Convex combination (cont.)

▶ This can be expressed equivalently as

$$\min_{\mathbf{w}_{T_i}, \xi_i, \lambda_i} \quad \frac{1}{2} \sum_{i=1}^{n} \frac{1}{\lambda_i} \|\mathbf{w}_{T_i}\|^2 + C \sum_{k=1}^{m} \xi_k$$

$$\text{s.t.} \quad \frac{1}{n} \sum_{i=1}^{n} \langle \mathbf{w}_{T_i}, \boldsymbol{\phi}_{T_t}(\mathbf{x}_k, \mathbf{y}_k) \rangle - \max_{\mathbf{y} \neq \mathbf{y}_k} \frac{1}{n} \sum_{i=1}^{n} \langle \mathbf{w}_{T_t}, \boldsymbol{\phi}_{T_i}(\mathbf{x}_k, \mathbf{y}) \rangle \geq 1 - \xi_k,$$

$$\xi_k \geq 0, \ \forall k \in \{1, \ldots, m\}, \ \sum_{i=1}^{n} \lambda_i = 1, \ \lambda_i \geq 0, \ \forall i \in \{1, \cdots, n\}.$$

Conclusions

- We show that if there is a learner $\mathbf{w}_G$ defined on a complete graph achieves a margin on some training data, then with a random collection of spanning tree learners $\{\mathbf{w}_{T_i}\}_{i=1}^n$ we can achieve a similar margin with high probability. Besides, n is polynomial in k.

- We propose two methods to combine the random collection of trees, namely, convex combination and conical combination.